

Nobody can draw your architecture. Your code already did.

The business and the engineering floor **stopped speaking the same language.**

On the ground, your engineers...

- Can't see what breaks downstream when they change a service
- Onboard in months, nobody can point to the whole system
- Rebuild what already exists on another team
- Now inherit AI-written code no one ever mapped

In the boardroom, you can't answer...

- What does our engineering org actually run?
- Which systems does this security issue touch?
- Where does sensitive data flow, for the audit?
- What's business-critical, and who owns it?

Three costs compounding. One lever that pays.



41%

of new code is AI-written



3_{mo}

to onboard an engineer
into an unmapped
codebase



0

real-time view leadership
has of what's actually
shipping



60%

fewer tokens for sharper
answers

A proven model, brought down to earth.

Palantir, for governments

- Builds an ontology of the organization's world
- Humans and AI reason over the same model
- Sits in the decision loop, not just a chart
- 18-month, eight-figure deployments

FOR


KnowledgeForge, for your software

- Builds an ontology of your codebase
- Engineers, execs and AI reason over the same map
- Generated from the source in minutes
- Self-serve, starts with one repo

— THE SOLUTION

We turn your repos into a living map.

Point it at a codebase, watch the C4 architecture graph build itself, straight from the source.



One graph. One language. No translation layer.



Your leadership

asks what we run, what's at risk, what this CVE touches, in plain business terms



Your engineers

ask what breaks if I change this, down to the file and the dependency



Your AI agents

pull curated context from the graph, sharper output, fewer tokens

Your agents are only as good as **the context you feed them.**

Context from the raw repo

- Greps blindly, dumps whole files
- Noisy and partial, often the wrong slice
- Weak on cross-repo links and business meaning
- More guessing, more hallucination
- And it burns tokens doing it

VS

Context from your graph

- The exact subgraph for the task, code and business
- Owners, dependencies, data flow, intent
- Curated structure → sharper, grounded answers
- Works on code AND business questions
- Side effect: ~60% fewer tokens, more AI on the same plan

Leadership gets a cockpit, not a status meeting.

Who's building what

activity and ownership per service, see where the effort actually goes

Where the risk is

compliance, security, stale or abandoned services, flagged on the map

The KPIs you choose

pin the metrics that matter to you, read straight off the C4 diagram

Plugs into your stack

connectors to the tools you already run, you steer what comes first

— WHAT IT'S WORTH TO YOU

The same work, a fraction of the time.

WHAT YOU DO TODAY	MANUAL WAY	WITH KNOWLEDGEFORGE	Δ
Audit & compliance prep	weeks	about a day	~10× faster
Onboard a new engineer	months	days	
Architecture diagram	hand-drawn, stale	auto, always current	
Blast radius of a new CVE	days of digging	one query	
AI agent context	raw repo dump	curated subgraph	sharper + ~60% fewer tokens

Point it once. Then it **runs itself**.



Point it at your repos

GitHub URL or local
folder, public or
private, one repo or
your whole estate



It builds the graph

parses the source
into a C4 ontology,
ecosystem,
services, internals,
with owners and
risk on every node



It stays current

re-extracts on
every commit, PR
and redeploy, the
map never drifts
from reality



Everyone queries it


people in plain
language, agents in
the same ontology,
one source of truth

— PROOF

Pointed at a real codebase, cold.

- A large multi-service codebase, never seen before
- Full C4 map, systems, services, dependencies, in **one pass**
- Every node: owners, tech stack, risk, compliance
- *Ask any node* in plain language

✓ Zero to a queryable map, nobody drew a thing.



See your own map. Heads you win, tails you lose nothing.

Send us one repo, we map it free and walk you through it on a 20-minute call.
No commitment, your data stays yours. A few founding-customer slots open this
quarter.

 [KNOWLEDGE-FORGE.COM](https://knowledge-forge.com) • IULIA@KNOWLEDGE-FORGE.COM • VLAD@KNOWLEDGE-FORGE.COM